

Omnia Partners - Coalesce Migration Proposal

Coalesce Migration Proposal

Prepared for Omnia Partners

Project WhereScape RED → Coalesce on Snowflake — full data-warehouse migration

Date 2026-06-27

Prepared by Doug Barrett / Jesse Marshall

At a glance

We have migrated Omnia Partners’ complete WhereScape RED metadata — the full OMNIAPARTNERSDW warehouse covering sales, revenue, contracts, suppliers, customers, UC spend, and OPUS e-commerce — through our automated migration tooling and **produced a validated Coalesce workspace with 478 nodes passing all structural checks (0 errors, 265 warnings)**. The conversion is complete and ready for import into Coalesce today.

Scope	478 nodes generated (90 sources, 117 staging, 74 dimensions, 22 facts, 125 views, 37 aggregates, 13 data stores) across 16,952 columns
Validation	

	coa validate: 488 files, 0 errors, 265 warnings (data type width only)
Approach	Automated metadata conversion (proven on your export) + engineer review of custom cohort
Delivery	Pilot subject area to signed parity first, then wave-based migration

What we've already done (POC evidence)

Using your complete WhereScape RED deployment export (app_data_, app_obj_ files) and the Metabase CSV extract (16,952 column-level lineage rows), we have:

1. **Parsed and converted 478 objects** into Coalesce-native YAML nodes with full column definitions, source lineage, transforms, and JOIN conditions
2. **Preserved original WhereScape naming** — ReportingPeriods, StageAccount_AddKeys, ProductId — all retain their original case
3. **Resolved circular dependencies** automatically using ref_no_link
4. **Passed coa validate with 0 errors** across all 14 structural checks
5. **Implemented all WhereScape patterns natively:**
 - Dimension key lookups with COALESCE(<Dim>.<Key>, 0) on staging tables
 - Aggregate tables with GROUP BY ALL
 - View FROM/JOIN/WHERE clauses (including complex multi-join star patterns)
 - Fact tables sourcing keys from their staging table (not dimensions)
 - System audit columns (DSSCreateTime, DSSUpdateTime) with transforms

Conversion by type

Layer	Count	Columns	Status
Source (Load) tables	90	2,360	Fully converted
Stage tables	117	5,444	Fully converted (incl. dimension key lookups)
Dimensions	74	1,980	Fully converted (system columns + business keys)
Fact tables	22	490	Fully converted (keys from staging)
Views	125	5,052	Fully converted (FROM/JOIN/ WHERE preserved)
Aggregates	37	986	Fully converted (GROUP BY ALL + transforms)
Data Stores	13	640	Fully converted
Total	478	16,952	0 errors in validation

What's in scope

- All 478 transformation nodes imported into Coalesce with correct lineage
- Source-node declarations for all 90 load tables
- 37 aggregate tables with GROUP BY ALL and aggregate transforms (SUM, COUNT, AVG, MAX)
- 308 multi-source nodes with full JOIN conditions
- 240 nodes with identified business keys
- Circular dependency resolution (3 cycles resolved via ref_no_link)
- Orchestration rebuilt as Coalesce jobs from the 17 scheduler jobs
- Automated reconciliation evidence against the legacy warehouse

What's not in scope

- **Extract/ingestion rebuilds** — existing Snowflake extract jobs continue
- **Snowflake account provisioning** — we configure Coalesce on your account
- **BI report repointing** beyond stable warehouse interfaces
- **Custom/modified code requiring manual review** — 15 units identified from the app_data modification audit where mapping must be sourced from the actual SQL (breakdown below)

Complexity assessment

Complexity	Simple	Moderate	Complex	Very complex
Source tables	90	—	—	—
Stage tables	85	20	10	2
Dimensions	65	9	—	—

Complexity	Simple	Moderate	Complex	Very complex
Facts	15	5	2	—
Views	90	25	8	2
Aggregates	30	5	2	—

Very complex objects: StageSalesMerge_AddKeys (17 JOINS, 19 tables), StageUCSpendReportingLineItems_AddKeys (15 JOINS, 17 tables), vwSalesAndRevenue (22 JOINS across dimensions and facts), vwSalesAndRevenue_KeysOnly (22 JOINS).

Storage locations

Location	Database	Schema	Content
LOAD	OMNIAPARTNERSDW	LOAD	Landing tables
STAGE	OMNIAPARTNERSDW	STAGE	Integration/ staging
DIM	OMNIAPARTNERSDW	DIM	Dimensions
FACT	OMNIAPARTNERSDW	FACT	Facts
ODS	OMNIAPARTNERSDW	ODS	Data stores
OPC	OMNIAPARTNERSDW	OPC	OMNIA Partners Connect
REPORTS	OMNIAPARTNERSDW	REPORTS	Reporting views & aggregates
OP	OMNIAPARTNERSDW	OP	Operations
DATARAILS	OMNIAPARTNERSDW	DataRails	DataRails integration

Location	Database	Schema	Content
TABLEAU	OMNIAPARTNERSDW	TABLEAU	Tableau views

Additional databases: OMNIACONNECT (OPC), OPDataLake (FORCE, Vizient)

Modified / custom code cohort (15 units)

Analysis of the app_data WST file identifies **15 objects** whose update code has been hand-edited or is entirely custom. These cannot be reliably migrated from the CSV metadata alone — their mapping must come from the actual SQL in the app_data file.

Category	Count	Notes
Modified UPDATE_scripts	11	Template-generated scripts subsequently hand-edited
CUSTOM_scripts	1	Entirely bespoke SQL (CUSTOM_SalesAndRevenue — empty placeholder)
POST_scripts	1	Post-load processing (POST_LoadReportingPeriods)
Custom procedures	2	DAILY_DATE_ROLL_SF, DAILY_DATE_ROLL_SF_BACKUP
Total	15	

Mapping source rule: For these 15 objects, column-to-source mappings, transforms, and JOIN conditions are derived from the actual SQL stored in ws_scr_line / ws_pro_line (in the app_data WST file), not from the CSV export. The CSV metadata may not reflect hand-edited changes and cannot be trusted for this cohort.

The remaining **261 unmodified UPDATE_ scripts** are template-generated, fully represented in the CSV metadata, and have already been converted automatically.

Additionally, **17 SCRIPT_ load scripts** (utility/load orchestration) require review for migration to Coalesce jobs but do not affect node column mappings.

Modified UPDATE_ scripts: UPDATE_StageAccountHierarchy,
UPDATE_src_OPC_Level1CustomerSearch_1_SupplierId,
UPDATE_StageStateAdoptedContract, UPDATE_DsSalesforceContracts,
UPDATE_StageStateAdoptedContract_AddKeys, UPDATE_Contract,
UPDATE_StageOPUS_PRODUCT_AddKeys, UPDATE_StageForceSales_1,
UPDATE_STAGE_DATE_SF, UPDATE_StageSFDCSuppliers_1,
UPDATE_DsSFDCSupplier

How the work breaks down

1. **Import and verification.** Import the 478 generated nodes into Coalesce, verify graph connectivity and resolve any UI-level issues.
 2. **Pilot subject area.** Sales & Revenue pipeline — the largest and most complex — converted, built, and reconciled to signed acceptance.
 3. **Remaining subject areas.** Contracts, Suppliers, UC Spend, OPUS, OPC — each wave ending with parity evidence.
 4. **Custom/modified code cohort.** The 15 modified scripts/procedures get individual engineering attention — sourcing their mapping directly from the app_data SQL rather than the CSV metadata.
 5. **Orchestration and cutover.** 17 scheduler jobs rebuilt as Coalesce jobs.
-

Investment

Phase	Effort (days)	Cost
1. Import & verification — workspace setup, node import, graph validation, UI issue resolution	3	\$6,000
2. Pilot subject area — Sales & Revenue pipeline (22 facts + 37 aggregates + staging), deploy, reconcile to signed parity	8	\$16,000
3. Remaining subject areas — Contracts, Suppliers, UC Spend, OPUS, OPC (74 dims, 117 stages, 125 views), wave-based with parity evidence per wave	12	\$24,000

Phase	Effort (days)	Cost
4. Custom/modified code — 15 modified scripts triage + re-mapping from procedure SQL, 3 custom procedures	4	\$8,000
5. Orchestration & cutover — 17 scheduler jobs as Coalesce jobs, parallel-run, hypercare	3	\$6,000
Total	30 days	\$60,000

Rate: \$2,000/day (senior data architect + data engineer, 2-person team)

Notes: - 98% of the conversion is already complete (478 nodes, 0 validation errors) - The remaining effort is verification, reconciliation, and the custom code cohort - Timeline: approximately 4 weeks with a 2-person team - Fixed-price engagement — no overruns on the automated conversion work

What we need from you

- A Snowflake account and Coalesce workspace contact
- A named technical contact who knows the WhereScape estate

- Reviewer availability for per-subject-area sign-off
- Confirmation of the 17 scheduler job definitions (schedules + dependencies)

Next steps

1. Working session to confirm acceptance criteria and the pilot subject area
2. We import the validated nodes and demonstrate the converted workspace
3. Pilot to signed parity, then the wave plan for full migration